

Linux Internals



ECE 373

Overview

- Short History
- Brief description
- 20,000 foot level
- Navigating the source
- Top Down trace – user land to kernel
- Bottom Up trace – hardware to kernel



Linux, at a glance

- General and flexible OS kernel
- Supercomputers, Servers, Clients, Devices
- Configurable
- Open Source
- GPLv2 (controversy!!)
- Usually found in a full distribution with user environments and GUIs

Pre-History

- UNIX – 1969, Thompson & Ritchie, PDP-11
- BSD - 1977, Bill Joy and friends, PDP and VAX
- GNU - 1983, Richard Stallman, et al
- MINIX - 1987, Tannenbaum, 80286
- Linux – 1991, Torvalds, 80386



Linux Early History

- 1990 – Started as a terminal emulator project
- 1991 – Released to public
- 1993 – Slackware distro, Peter Volkerding
- 2000 – RedHat Commercial distro
- 1999 – Realtime extensions
- Embedded uses – Wind River, MontaVista, etc

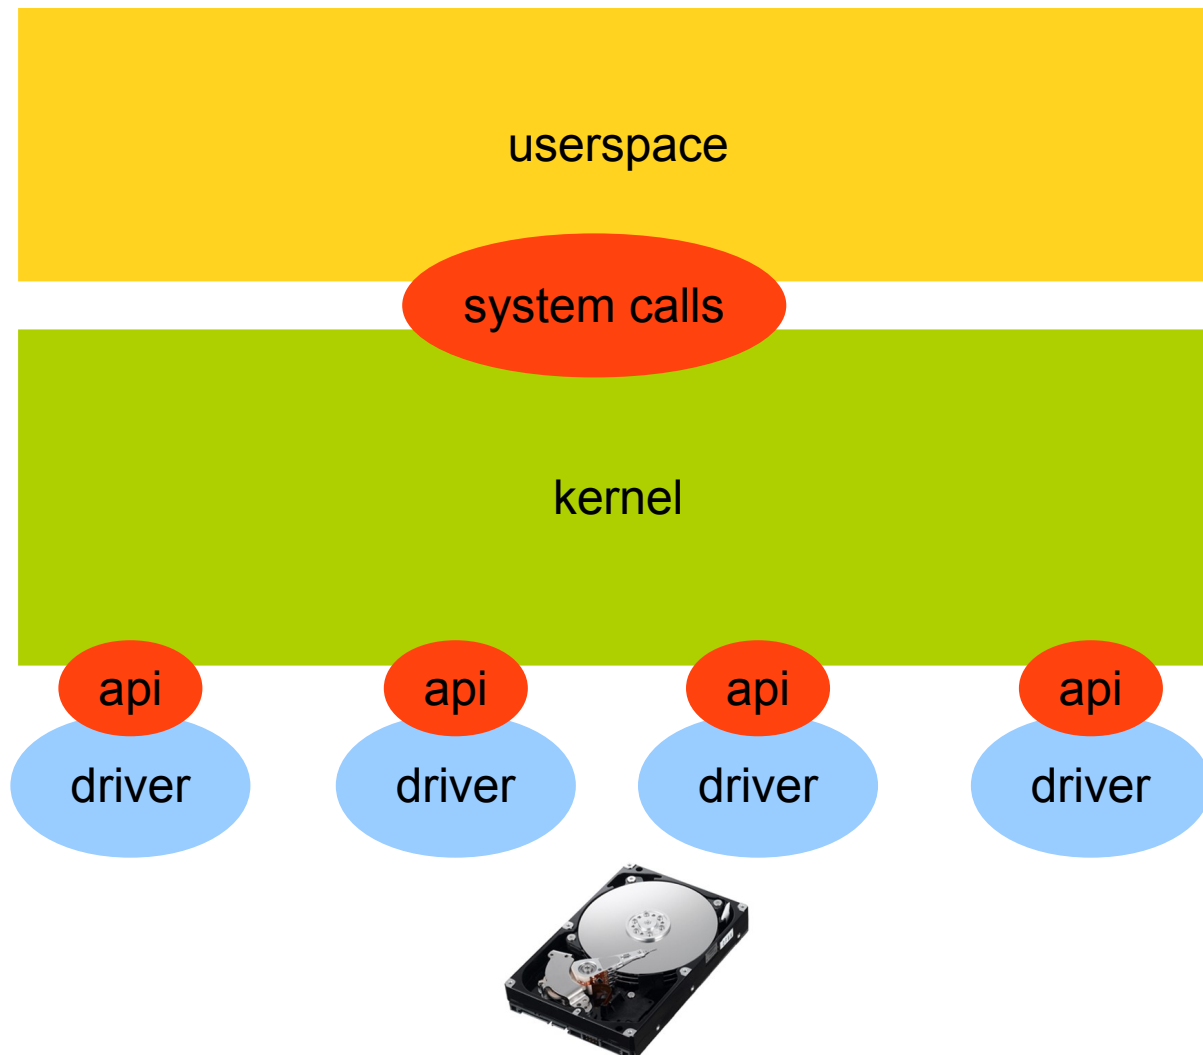


20,000 foot level

- Monolithic kernel
- Dynamically loadable modules
- Pre-emption

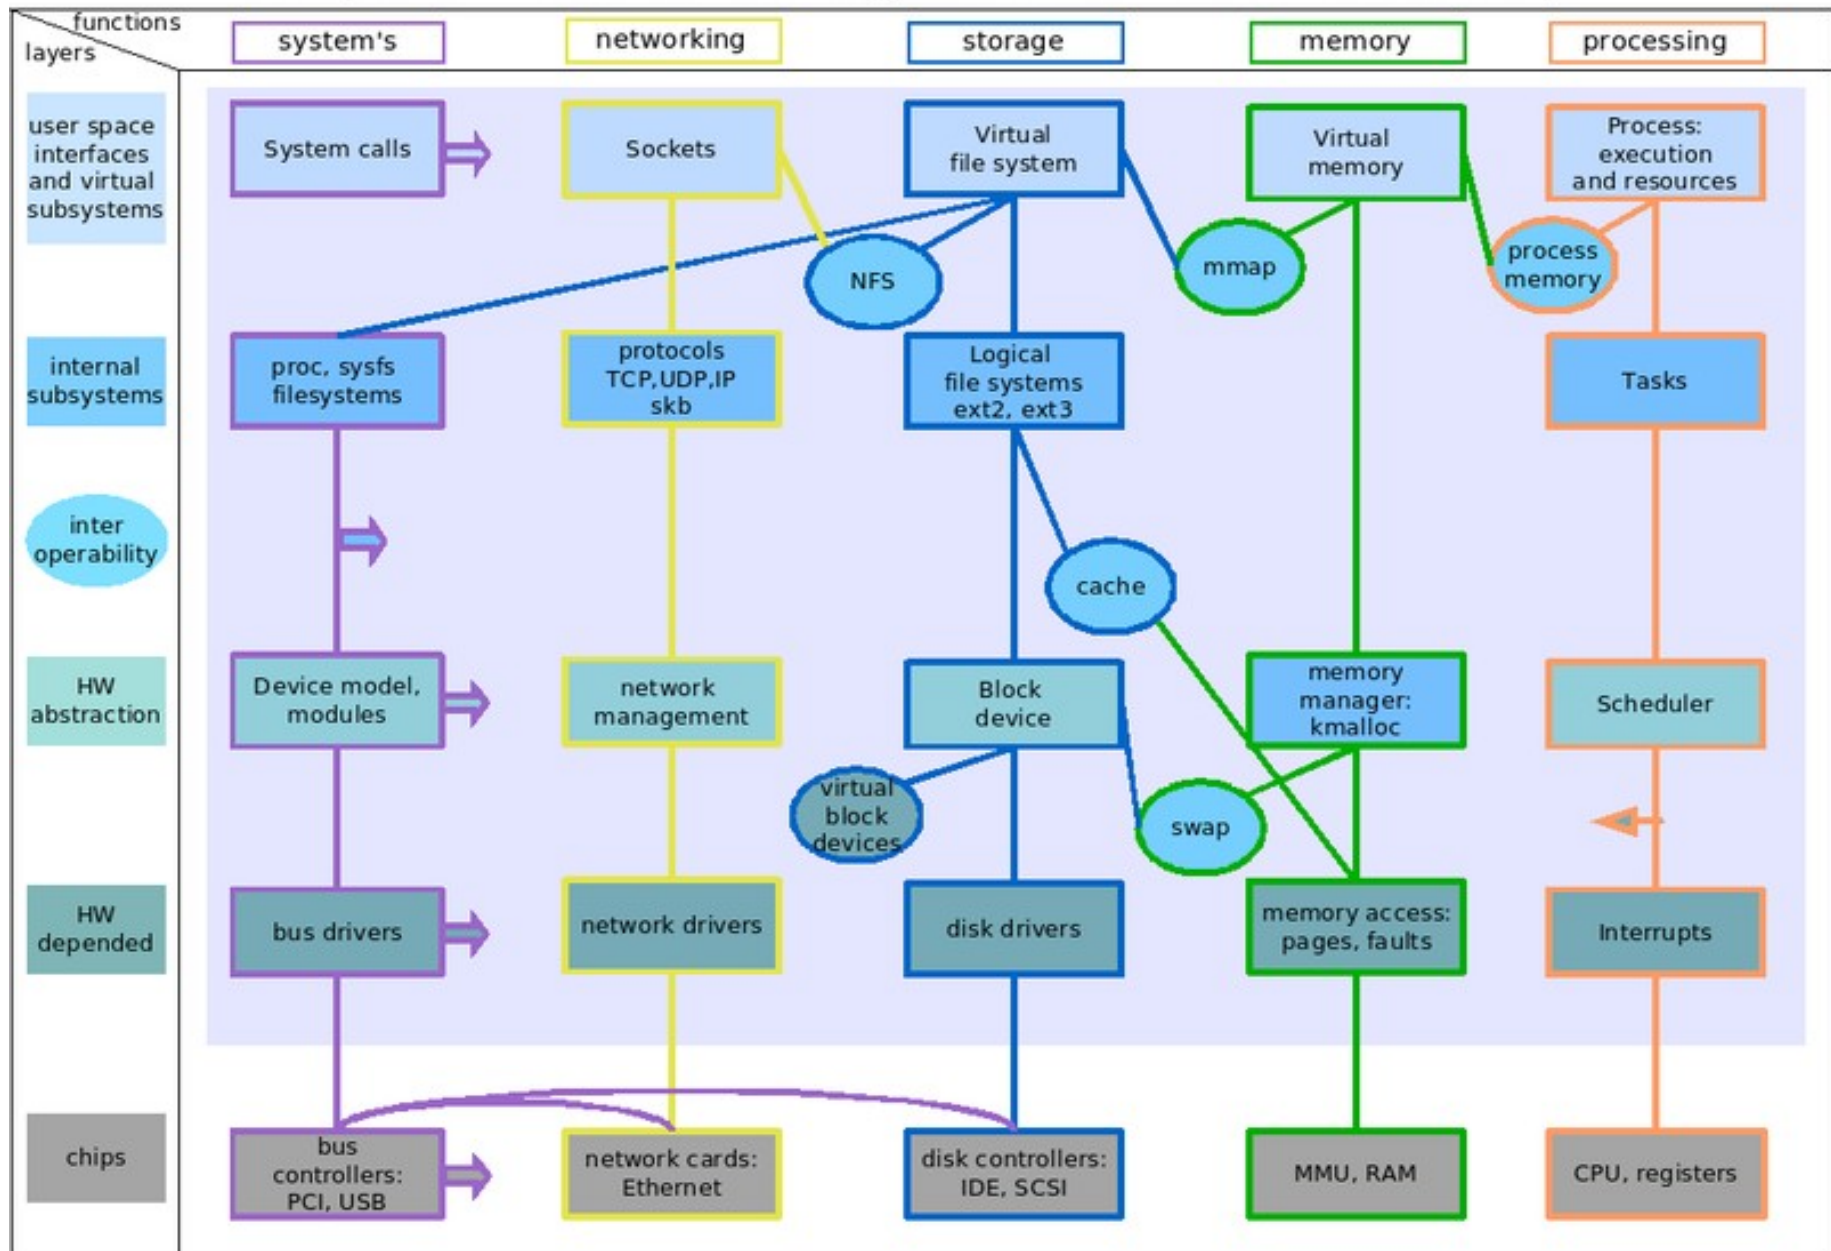


A little closer...



Too close?

Simplified Linux kernel diagram in form of a matrix map



Navigating the dark and creepy places


- Where to start when diving in?
- Many resources to read kernel
- kernel.org + text editor
 - <https://www.kernel.org/>
- ctags / cscope
- LXR resources
 - <http://lxr.free-electrons.com/>



Other Kernel Resources

http://elinux.org/Linux_Kernel_Resources


- References
 - kernel source
 - mailing lists
 - books
 - news
 - alternate architectures



Top down

What happens when a program requests the current time?

- program calls gettimeofday()
- libc
- Systrap with user buffer
- Kernel finds RTC driver
- RTC driver gets time value from clock HW
- Kernel service writes time to user space
- Systrap return



Following gettimeofday()

- First look at the manpage!
- Call starts in libc (not within today's scope)
- gettimeofday() wrapper called in kernel
- System call initiated, trap invoked through fast software interrupt
- Clock source read
- Buffers returned
- Trap complete



Walking the path...

- <http://elixir.free-electrons.com/linux/latest/source/include/uapi/asm-generic/unistd.h#L483>
- <http://lxr.free-electrons.com/source/include/linux/syscalls.h#L212>
- <http://lxr.free-electrons.com/source/arch/arc/kernel/sys.c#L10>
- <http://lxr.free-electrons.com/source/kernel/time/time.c#L102>
- <http://lxr.free-electrons.com/source/kernel/time/timekeeping.c#L695>
- <http://lxr.free-electrons.com/source/kernel/time/timekeeping.c#L526>
- <http://lxr.free-electrons.com/source/kernel/time/timekeeping.c#L493>
- <http://lxr.free-electrons.com/source/kernel/time/timekeeping.c#L194>
- ... to a registered clock driver
- http://en.wikipedia.org/wiki/System_call



Bottom up

- What happens when the user hits "<enter>"?
 - Serial line read by UART
 - UART buffers characters, pulls interrupt line #x
 - CPU stores current process A, searches table for handler for interrupt #x
 - Handler code extracts character from UART through I/O register access
 - Find process B waiting on device semaphore, give it the character
 - Reschedule process B as Ready
 - Return from interrupt and resume process A

This was only an overview...

- Linux is open, readable, but big
- Tons of resources available to traverse the kernel and its code
- Will drill much deeper into the kernel as class goes on, don't worry!
- Ask questions!!

