

DMA and Descriptors



ECE 373

Class notes for today

- DMA – what is it, why do we care?
- How DMA works
- What existed before DMA?
- How DMA is setup in Linux
- Descriptors, huh?
- Bonus: Packing command buffers

K*alloc and friends

- Many ways to alloc heap memory
 - `kmalloc()`, `kmalloc_array()`, `kmalloc_node()`
 - `kzalloc()`, `kzalloc_node()`,
 - more....
- `void *kmalloc(size_t size, gfp_t flags)`
 - Flags: `GFP_USER`, `GFP_KERNEL`, `GFP_ATOMIC`, `GFP_NOWAIT`, `GFP_DMA`, ...
- `void kfree(void *)`
- See kernel source


<http://lxr.free-electrons.com/source/include/linux/slab.h#L359>

Direct Memory Access, the story

- DMA is a feature found in many pieces of hardware
- Allows direct access to system memory without the CPU
- Many devices live on slower busses
- DMA allows slower devices to transfer data at their leisure


Why do we care about DMA?

- If devices can't move their own data, who does?
- CPU can only do one thing at a time
- If CPU is pre-occupied with copying data, system becomes slooooow



Back in the day...

- DMA didn't always exist
- Device RAM mapped into system memory
- CPU would do all the work moving data around
- Disk controllers worst, since disk spindles very slow compared to main memory
- If the CPU waits for data, it will stall
- Stall = Bad

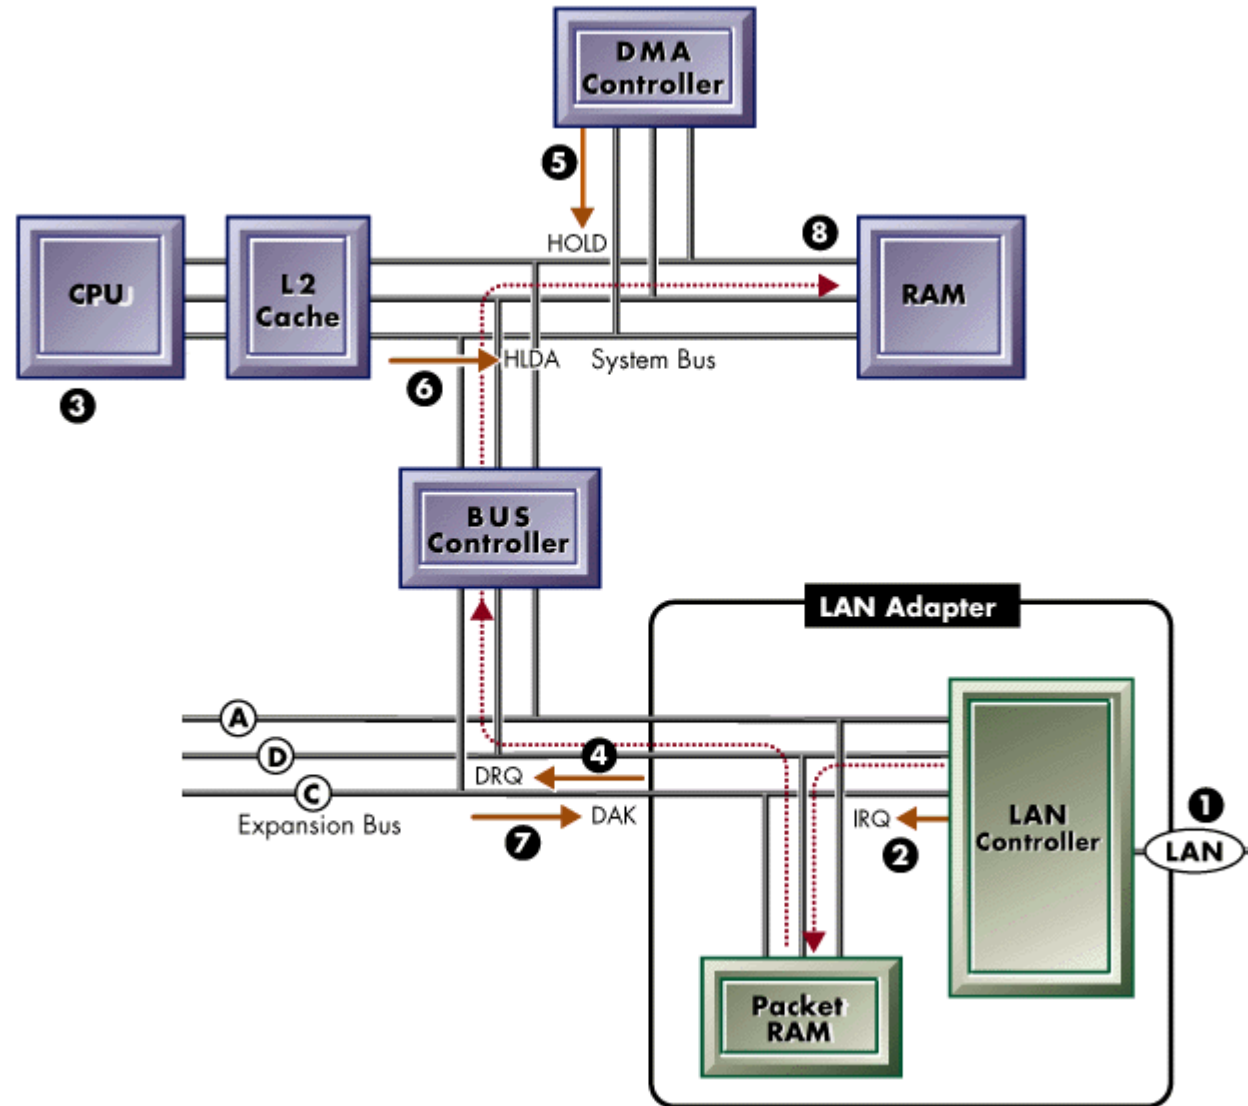


Evolution of DMA

- DMA-capable devices showing up on the ISA bus
- Used a central DMA controller in between peripheral devices and memory
- DMA controller acted like a proxy
- Devices still using ISA DMA today are floppy disks and old parallel ports

Controller Driven DMA

- 1 – Data received from network
- 2 – Device sends IRQ
- 3 – IRQ handler starts DMA
- 4 – Device requests DMA access
- 5 – DMA controller requests a bus hold
- 6 – CPU ack's the hold
- 7 – Bus controller ack's the DMA
- 8 – Device copies data to RAM



Further evolution of DMA

- DMA in PCI and PCI Express much different than ISA
- No centralized DMA controller
- DMA is performed by peripheral device
- However, not all PCI devices allowed to talk at once

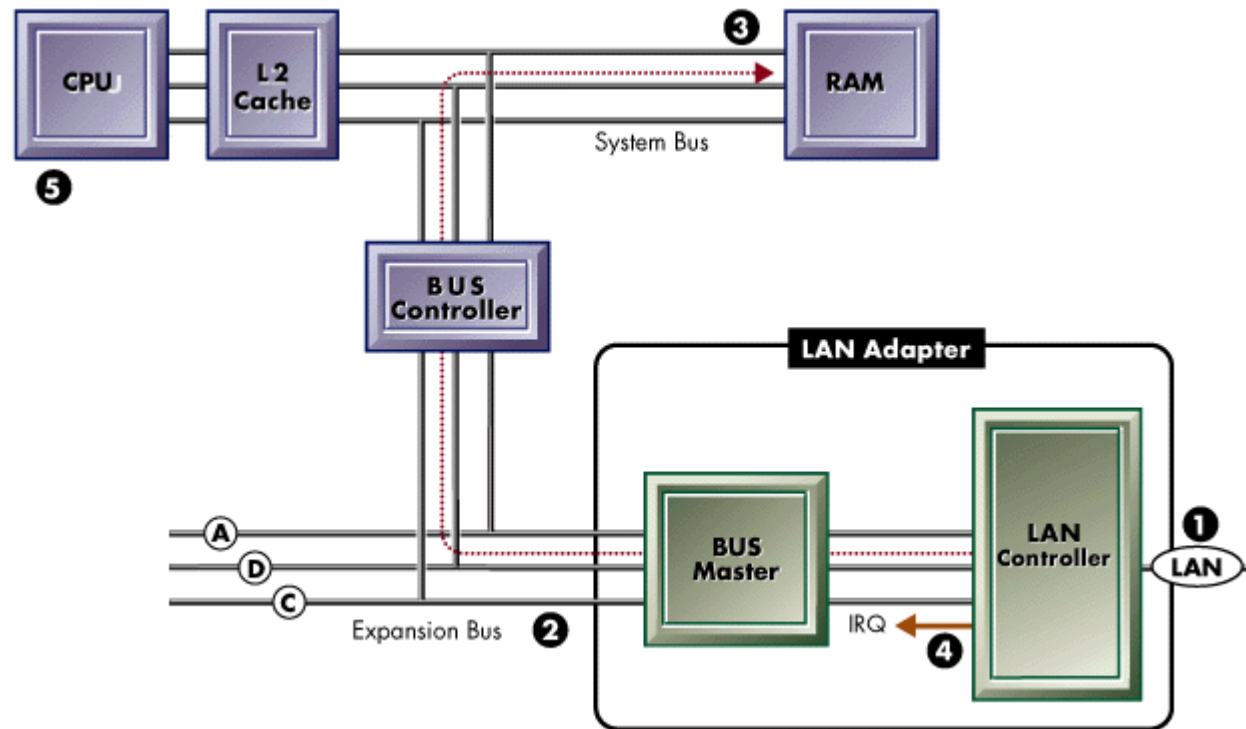
There can be only one...

- PCI devices utilize protocol called bus mastering
- Each device requests ownership of PCI
- PCI controller arbitrates over all bus masters
- Each bus master issues commands when he owns the bus



Device Driven DMA

- 1 – Data received from network
- 2 – Bus master negotiates for bus access
- 3 – Device copies data to RAM
- 4 – Device sends IRQ
- 5 – IRQ handler processes data



DMA, how it works

- DMA-enabled device told which direction to DMA (to or from device, or bidirectional)
- Physical memory is mapped for DMA
- Device initiates the DMA (either to or from device)
- Memory unmapped for DMA after operation completes

DMA in Linux

- Function APIs exist for each driver type to control DMA
- DMA consists of mapping memory for DMA, unmapping when finished
- Mapping memory means pinning it down, not allowing it to be swapped out by memory manager
- DMA deals with physical address (or bus address)

DMA in PCI

- API for allocating and mapping memory in PCI:
 - `dma_alloc_coherent(dev, size, dma_handle, flags)`
- Coherent means can safely be used for both IN and OUT
- DMA handle returns (by reference) physical address of mapped memory
- Return type of function is kernel logical address of mapped memory
- Flags dictate mem allocation details: `GFP_ATOMIC`, `GFP_KERNEL`, `__GFP_DMA`

Contrived Example

```
/*
 * contrived DMA example
 */
static void ece_move_my_data(struct dev mydev, int len,
                             unsigned char *buf, void *hw_addr)
{
    dma_addr_t daddr;
    unsigned int reg;
    unsigned char *pmem;

    /* get and prepare the DMA memory */
    pmem = dma_alloc_coherent(mydev, len, &daddr, GFP_KERNEL);
    memcpy(pmem, buf, len);

    /* tell the device where the DMA space is */
    reg = (daddr >> 32) & 0xffffffff;
    writel(reg, hw_addr + ECE_DMA_HI_REGISTER);
    reg = daddr & 0xffffffff;
    writel(reg, hw_addr + ECE_DMA_LO_REGISTER);

    /* tell the device to copy the data */
    reg = 1;
    writel(reg, hw_addr + ECE_READ_DOORBELL);

    /* wait for DMA to complete */
    /* an interrupt handler would be better */
    usleep(10);
    dma_free_coherent(mydev, len, daddr);
}
```

Mapping existing memory

- `kmalloc()` returns kernel logical address
- Kernel logical addresses can be pinned
- `dma_map_single(dev, ptr, size, flags)` maps existing kernel logical address to pinned physical address
- Handy for mapping only when needed, keeps pinned memory space small
- Flags dictates copy direction, either `DMA_TO_DEVICE`, `DMA_FROM_DEVICE`, or `DMA_BIDIRECTIONAL`

Another Contrived Example

```
/*
 * contrived DMA example
 */
static void ece_move_my_data(struct dev mydev, int len,
                             unsigned char *buf, void *hw_addr)
{
    dma_addr_t daddr;
    unsigned int reg;
    unsigned char *pmem;

    /* pin the existing buffer */
    daddr = dma_map_single(mydev, buf, len, DMA_TO_DEVICE);

    /* tell the device where the DMA space is */
    reg = (daddr >> 32) & 0xffffffff;
    writel(reg, hw_addr + ECE_DMA_HI_REGISTER);
    reg = daddr & 0xffffffff;
    writel(reg, hw_addr + ECE_DMA_LO_REGISTER);

    /* tell the device to copy the data */
    reg = 1;
    writel(reg, hw_addr + ECE_READ_DOORBELL);

    /* wait for DMA to complete */
    /* an interrupt handler would be better */
    usleep(10);
    dma_unmap_single(mydev, daddr, len, DMA_TO_DEVICE);
}
```

DMA engines for non-devices

- Other areas of an OS require memory copies
- TCP stack needs to copy kernel memory to user memory (think `copy_to_user()`)
- Intel I/O AT was DMA engine in chipset
- Changes made to OS's to offload these copies to I/O AT silicon
- Offloaded those particular memcpy operations from main CPU

Managing lots of DMA

- Most things that do DMA need to do lots of it
 - Disk controllers writing to and reading from disks
 - Network controllers receiving and transmitting
- We use lists of DMA requests
- Each request describes the activity
 - What data to move
 - How much data
 - In or Out
- These are called "descriptors"

Descriptors

- What is a descriptor?
 - Hardware field describing what work to do
 - Hardware field describing what work was done
- Carries bits and fields as instructions
- Carries pointers to buffers needing to be DMA'd into or out of hardware
- May return result information after the action

Example Descriptor layout

Table 7-24. Transmit Descriptor (TDESC) Fetch Layout - Legacy Mode

	63 48	47 40	39 36	35 32	31 24	23 16	15 0
0	Buffer Address [63:0]						
8	VLAN	CSS	Reserved	STA	CMD	CSO	Length

Table 7-25. Transmit Descriptor (TDESC) Write-Back Layout - Legacy Mode

	63 48	47 40	39 36	35 32	31 24	23 16	15 0
0	Reserved				Reserved		
8	VLAN	CSS	Reserved	STA	CMD	CSO	Length

Note: For frames that span multiple descriptors, the *VLAN*, *CSS*, *CSO*, *CMD.VLE*, *CMD.IC*, and *CMD.IFCS* are valid only in the first descriptors and are ignored in the subsequent ones.

Descriptors in code, simple

```
/* Transmit Descriptor */
struct e1000_tx_desc {
    __le64 buffer_addr;    /* Address of the descriptor's data buffer */
    __le64 flags_fields;
};
```

... or ...

```
struct e1000_tx_desc {
    __le64 buffer_addr;    /* Address of the descriptor's data buffer */
    __le32 lower_fields;
    __le32 upper_flags;
};
```

... or useful

```
/* Transmit Descriptor */
struct e1000_tx_desc {
    __le64 buffer_addr;           /* Address of the descriptor's data buffer */
    union {
        __le32 data;
        struct {
            __le16 length;        /* Data buffer length */
            u8 cso;              /* Checksum offset */
            u8 cmd;              /* Descriptor control */
        } flags;
    } lower;
    union {
        __le32 data;
        struct {
            u8 status;           /* Descriptor status */
            u8 css;             /* Checksum start */
            __le16 special;
        } fields;
    } upper;
};
```

Filling it out

```
buffer_info->buf = databuf;  
buffer_info->length = len;  
buffer_info->dma = dma_map_single(mydev, databuf,  
                                len, DMA_TO_DEVICE);
```

```
tx_desc = E1000_TX_DESC(*tx_ring, i);  
tx_desc->buffer_addr = cpu_to_le64(buffer_info->dma);  
tx_desc->lower.flags.length = cpu_to_le16(buffer_info->length);  
tx_desc->lower.flags.cso = 0;  
tx_desc->lower.flags.cmd = E1000_TXCMD_EOP;  
tx_desc->upper.data = 0;  
tx_desc->upper.fields.special = cpu_to_le16(vlan_id);
```

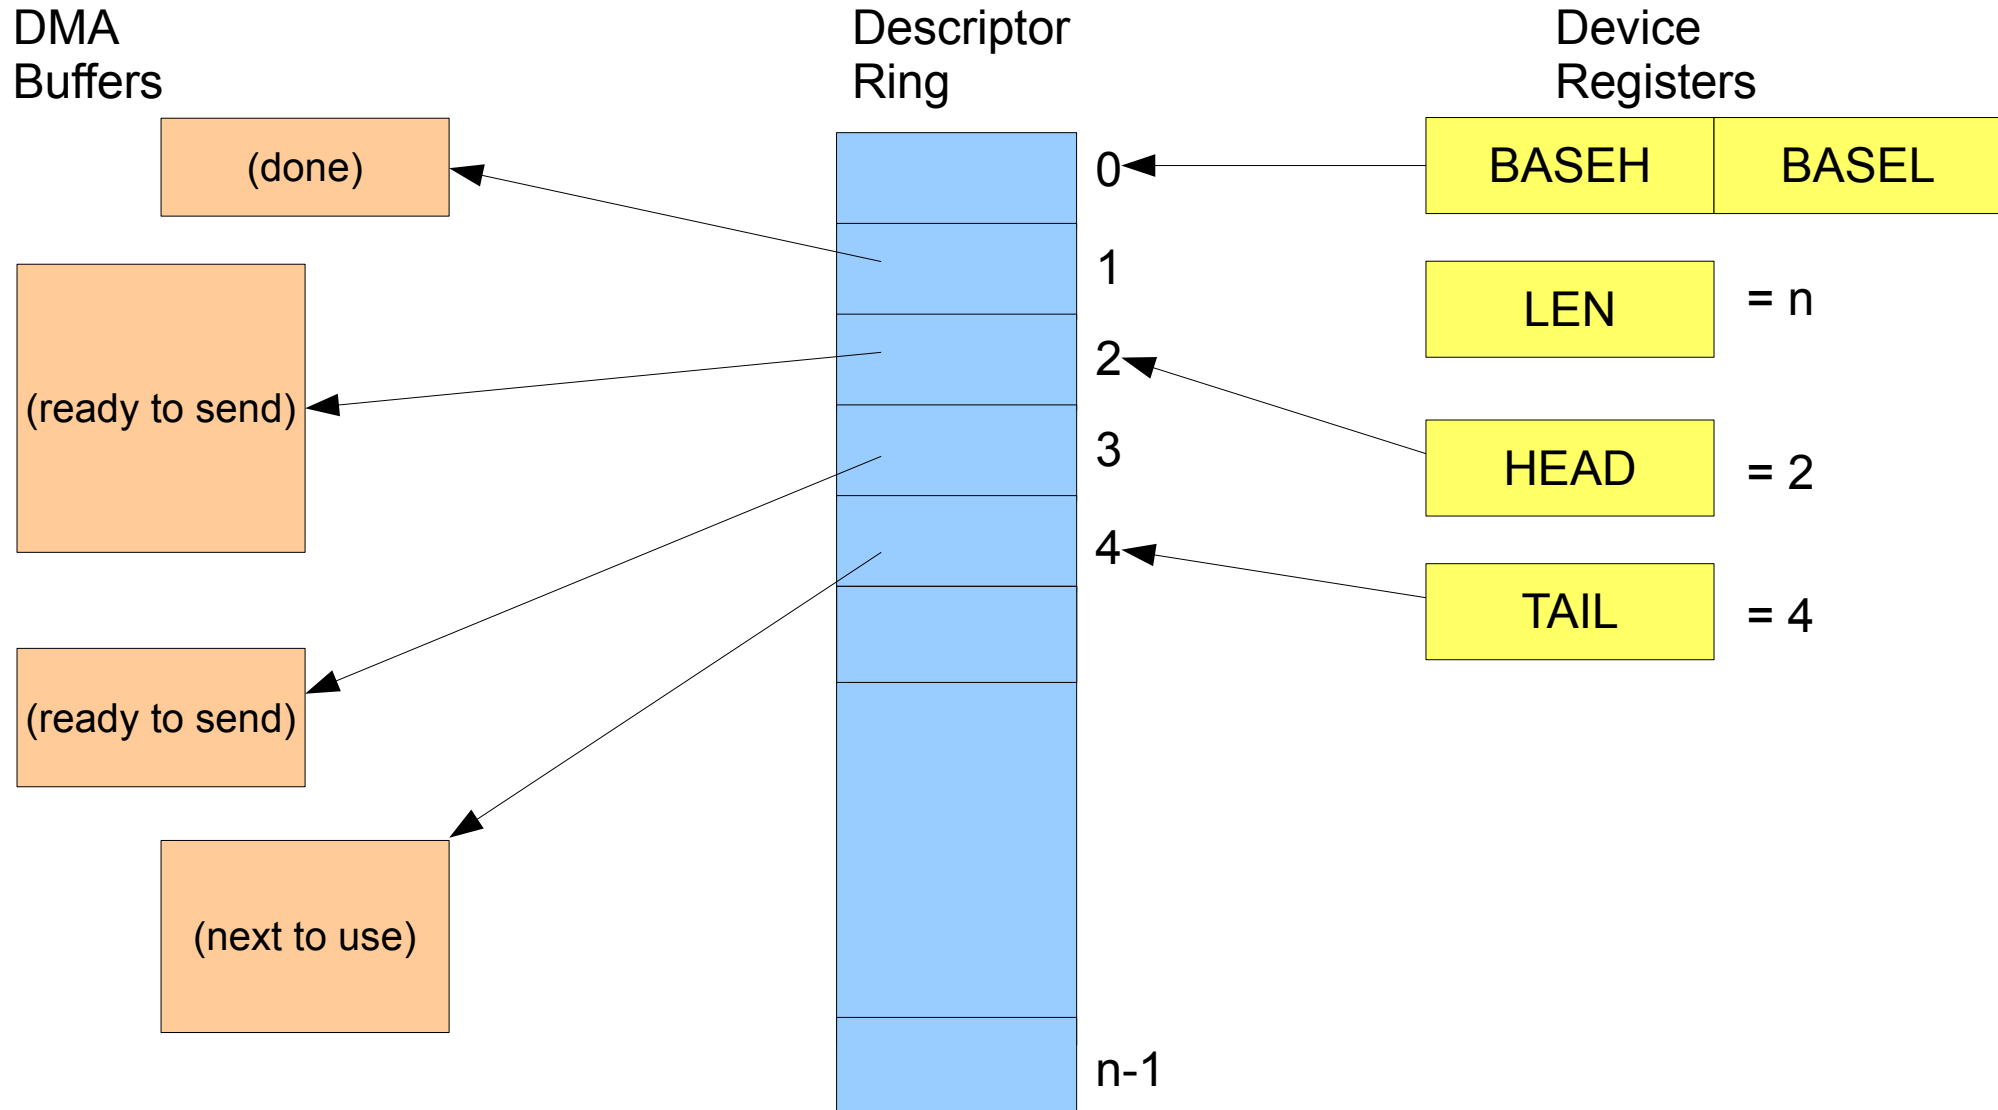

Rings!!

- Descriptors stored in ring buffer
 - Long array of descriptors
 - HW uses each through the array, then loops back
- Driver sets up each descriptor in ring as needed
- In order for hardware to use, what must we do?

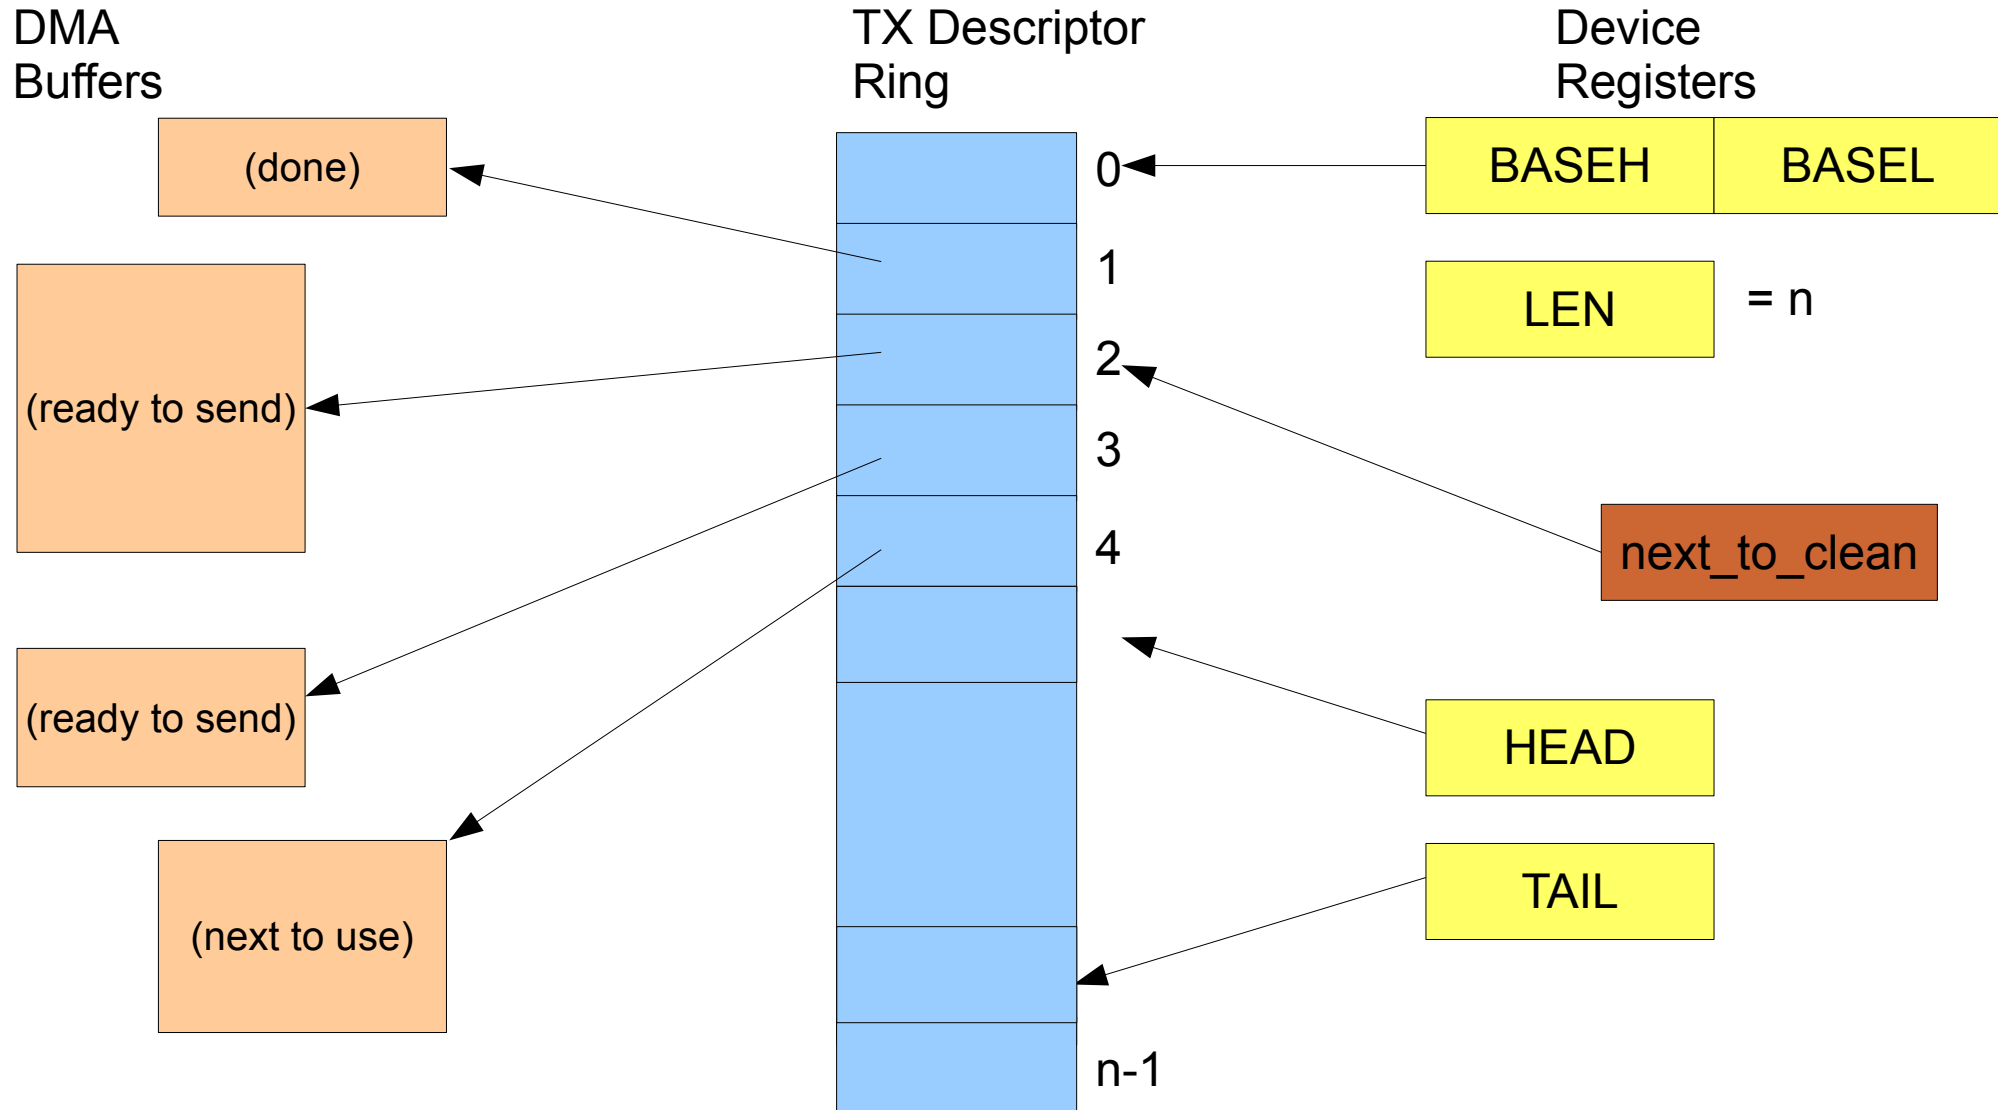
- Pin it!



Cheesy Picture



Another Cheesy Picture



Filling out a ring

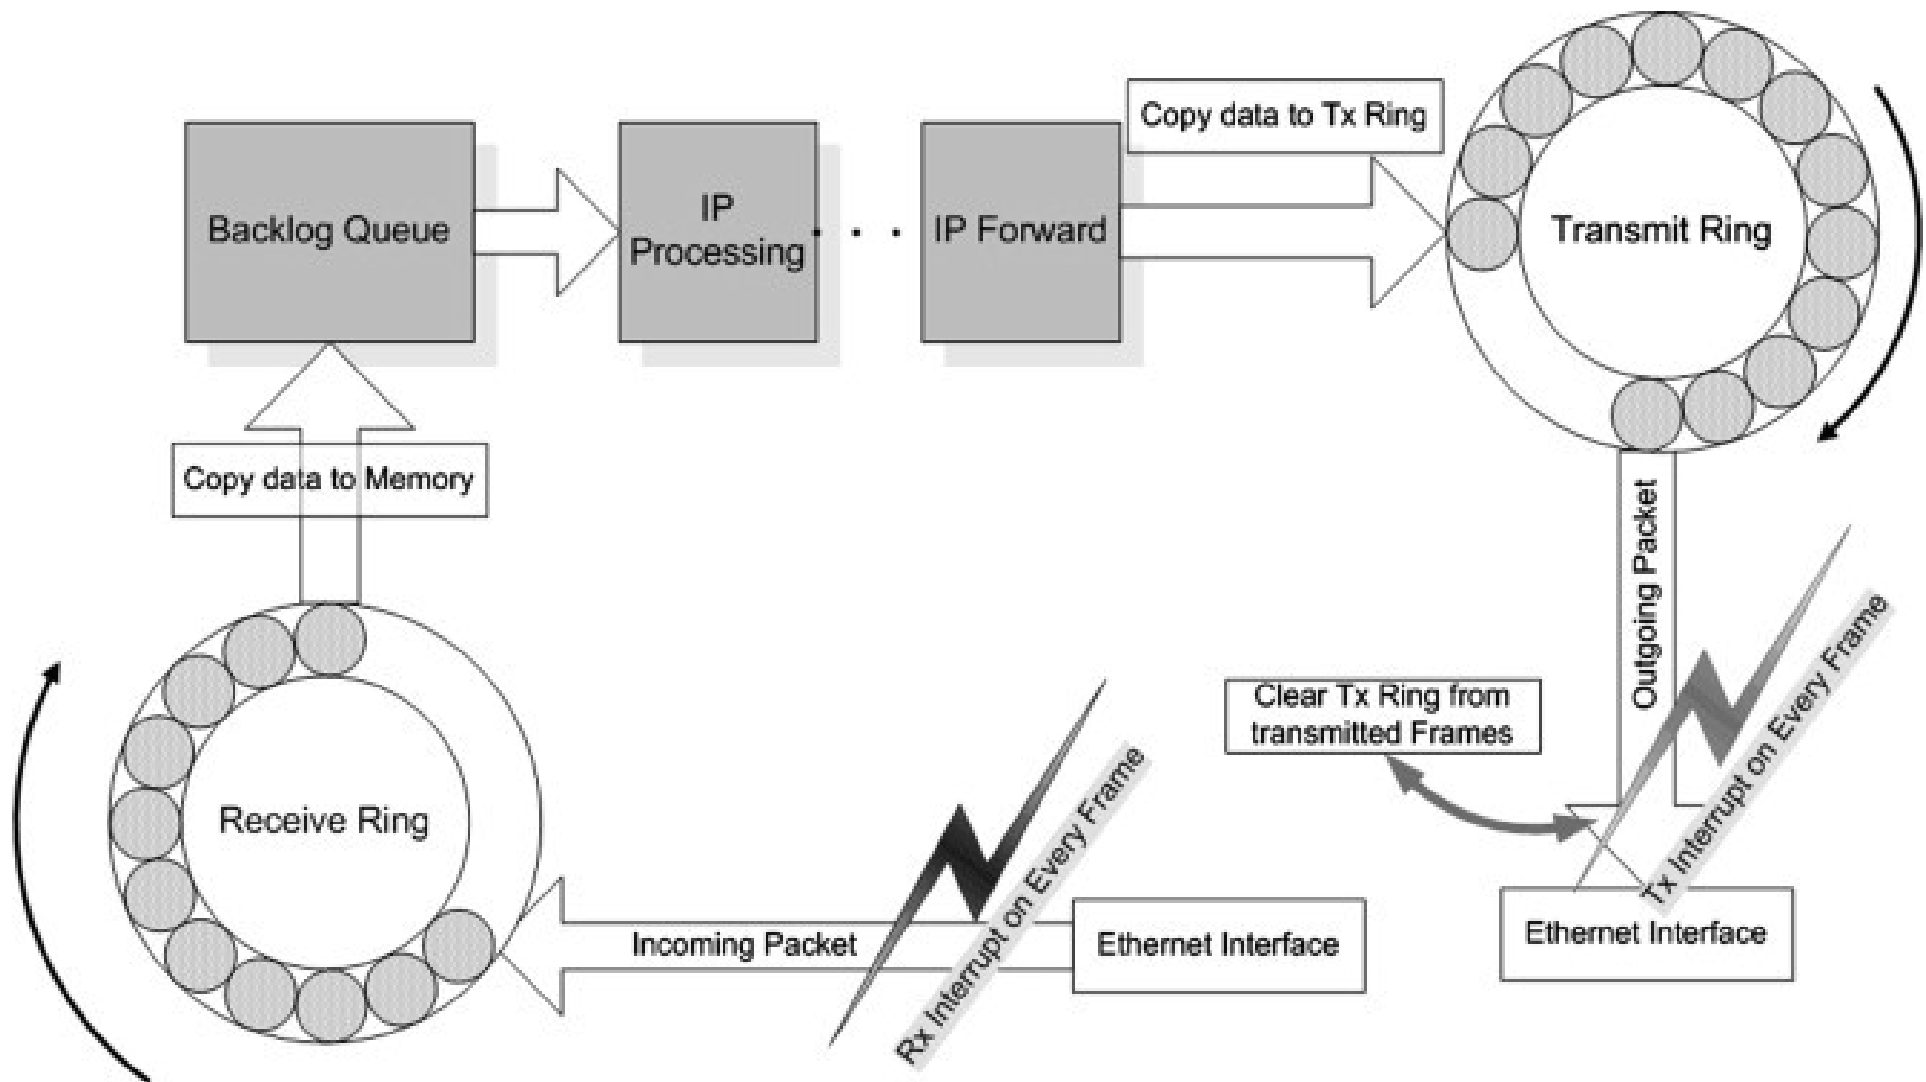
```
do {
    buffer_info = &tx_ring->buffer_info[i];
    tx_desc = E1000_TX_DESC(*tx_ring, i);
    tx_desc->buffer_addr = cpu_to_le64(buffer_info->dma);
    tx_desc->lower.data = cpu_to_le32(txd_lower | buffer_info->length);
    tx_desc->upper.data = cpu_to_le32(txd_upper);

    i++;
    if (i == tx_ring->count)
        i = 0;
} while (--count > 0);

tx_desc->lower.data |= cpu_to_le32(adapter->txd_cmd);

writel(i, E1000_TDT); /* bump tail */
```

Managing rings



Setting up a network device

- Reset device
- Create descriptors
- Set promiscuous mode
- Set autoneg speed
- Force link up
- Start interrupt handlers

Sniffing around

- Wireshark
 - Good frontend to packet tracing applications
 - Used for anything that can stream data
- pcap
 - Packet capture file
- tcpdump
 - Unix/Linux-based packet tracer
- Save as a standard libpcap!

Play the packet trace

- Linux: tcpreplay
- Windows: PlayCap
- Takes pcap from previous trace, plays it back on the network
- Good for repeatable testing

Sending test packets

- Single packet (linux)
 - `ping -I ethX -r 11.22.33.44`
 - Other arguments let you shape the packet
- Packet streams
 - Linux: `tcpreplay`, Windows: `PlayCap`
 - Plays back packet capture from previous trace (`tcpdump`)
 - Good for repeatable testing

Readings

- LLD3 pg 440-453
- ELDD pg 301-310
 - Uses pci_xxx() rather than dma_xxx() API
- <http://lxr.linux.no/#linux+v2.6.35.9/Documentation/DMA-API-HOWTO.txt>
- <http://lxr.linux.no/#linux+v2.6.35.9/Documentation/DMA-API.txt>