

Interrupts and Handlers




ECE 373

HW Interrupts

- Hardware wants attention
 - Data waiting, might be time-sensitive
- Interrupt handlers
 - Temporarily take over current thread, whether kernel or user
 - Can't be scheduled, can't sleep
 - no `msleep()`, no `mutex()`, careful with `kzalloc()`...
- Interrupts interrupting interrupts?



Thread, interrupted



Basic interrupt flow

- Allocate interrupt and set up handler
 - specific to what your driver needs to process
 - USB device would clean URBs, possibly refill if needed
 - Network device will process received packets
 - Disk device will process data blocks
- Tell device what interrupt line/number to use
- Run...
- Disable interrupt in device then OS when done

Interrupt Types

- Legacy
 - Hardware based, individual wires, limited availability
 - Chaining of handlers on same interrupt – coordination of handlers
- MSI – PCI message signaled interrupt
 - No more chaining,
 - Each PCI device gets its own interrupt
- MSI-X – MSI eXtended
 - Many interrupts per PCI device
- SW interrupts
 - SW triggers, Timers





Basic Wiring

- Request an IRQ using device info

```
err = request_irq(pdev->irq, ece_irq_handler, 0, "ece_int",  
data);
```

<http://lxr.free-electrons.com/source/include/linux/interrupt.h>

- irq: interrupt index in OS tables
- ece_irq_handler: ptr to interrupt handler function
- 0: flags
 - IRQF_SHARED, IRQF_SAMPLE_RANDOM, IRQF_TRIGGER_*
- "ece_int": name seen in last column of /proc/interrupts
- data: "magic cookie" data passed into intr handler
 - Same as in the timer handler

Legacy and MSI Setup

- Legacy

- Request an IRQ using PCI device info

```
err = request_irq(pdev->irq, ece_irq_handler, 0, "ece_int",  
data);
```

- MSI

- Enable MSI for the device
 - Request an IRQ using PCI device info

```
pci_enable_msi(pdev);
```

```
err = request_irq(pdev->irq, ece_irq_handler, 0, "ece_msi",  
data);
```

MSI-X Setup

MSI-X needs an array for many interrupt vectors

- Prep an array of msix data structures

```
entries = kcalloc(num_vectors, sizeof(struct  
msix_entry), GFP_KERNEL);
```

- Enable MSI-X for the device


... and get block of interrupt vectors

```
pci_enable_msix(pdev, entries, num_vectors)
```

- Request an IRQ for each msix entry

```
for (i = 0; i < num_vectors; i++)
```

```
err= request_irq(entries[i].vector,  
ece_irq_handler, 0, "ece_msix", i);
```



Magic cookie,
irq context

Example: MSI-X

```
struct msix_entry *msix_list;
char v_name[16];

msix_list = kcalloc(v_num, sizeof(struct msix_entry), GFP_KERNEL);
if (NULL == msix_list)
    return NULL;

/* prep the vector array */
for (v = 0; v < v_num; v++)
    msix_list[v].entry = v;

while (v_num >= least_vectors_needed) {

    /* try to get a block of vectors */
    err = pci_enable_msix(pdev, msix_list, v_num);

    if (0 == err)                /* success */
        break;
    else if (err < 0)            /* nasty failure, quit now */
        v_num = 0;
    else                         /* err == num vectors we should try */
        v_num = err;
}

/* failed, so clean up and return */
if (v_num < least_vectors_needed) {
    kfree(msix_list);
    msix_list = NULL;
    return NULL;
}

/* init all the vectors */
for (v = 0; v < v_num; v++) {
    snprintf(v_name, sizeof(v_name), "ece373_v_%02d", v);
    err = request_irq(msix_list[v].vector, ece_irq_handler,
                     0, v_name, ece_data);
}
ece_data->msix_list = msix_list;
```

Another example: MSI-X

```
/**
 * i40e_reserve_msix_vectors - Reserve MSI-X vectors in the kernel
 * @pf: board private structure
 * @vectors: the number of MSI-X vectors to request
 *
 * Returns the number of vectors reserved, or error
 */
static int i40e_reserve_msix_vectors(struct i40e_pf *pf, int vectors)
{
    vectors = pci_enable_msix_range(pf->pdev, pf->msix_entries,
                                   I40E_MIN_MSIX, vectors);
    if (vectors < 0) {
        dev_info(&pf->pdev->dev,
                 "MSI-X vector reservation failed: %d\n", vectors);
        vectors = 0;
    }

    return vectors;
}
```

```
    if (!adapter->msix_entries) {
        adapter->msix_entries = kcalloc(num_msix,
                                         sizeof(struct msix_entry),
                                         GFP_KERNEL);

        if (!adapter->msix_entries)
            return -ENOMEM;
    }

    for (vector = 0; vector < num_msix; vector++)
        adapter->msix_entries[vector].entry = vector;

restore:
    err = pci_enable_msix_exact(pdev, adapter->msix_entries, num_msix);
    if (err == -ENOSPC) {
        if (!adapter->drv_tss_rings && !adapter->drv_rss_rings)
            return err;

        netdev_info(adapter->netdev,
                     "Unable to allocate %d MSI-X vectors, Available vect
ors %d\n",
                     num_msix, err);
    }
}
```

Handler

- Interrupt handler is called with cookie argument for context
- Handler often needs to check interrupt cause register in device

```
int_bits = readl(hw_addr + INT_CAUSE)
```

- Returns
 - IRQ_NONE – not mine, shared by some other handler
 - IRQ_HANDLED – done and handled
 - http://lxr.free-electrons.com/source/drivers/net/ethernet/intel/i40e/i40e_main.c#L3070



Handler be quick!

- Blocking other interrupt handling and user jobs
- Grab HW info, stash away for later
- Don't call code that might sleep
 - sleep(), malloc(), other I/O functions
 - kmalloc GFP_ATOMIC maybe
 - Scheduler can't put interrupt handlers in wait queue
- Locks?
 - no mutex or semaphore
 - atomics and completions okay
- Use Top half / Bottom half concept
 - Wake up driver code with worker thread or waiting on a completion



Handler Code

```
#define REG_IRQ_CAUSE_READ_DONE    0x0001
#define REG_IRQ_CAUSE_WRITE_DONE   0x0002
#define REG_IRQ_CAUSE_ON_FIRE      0x0004

static irqreturn_t ece_irq_handler(int irq, void *data)
{
    struct ece_data_t *ece_data = data;
    u32 cause;

    /* no printing here - can't do anything that might sleep */
    irq_info = irq;

    cause = readl(hw_addr + REG_IRQ_CAUSE);

    switch (cause) {
    case REG_IRQ_CAUSE_READ_DONE:
    case REG_IRQ_CAUSE_WRITE_DONE:
        schedule_work(ece_data->io_task);
        break;
    case REG_IRQ_CAUSE_ON_FIRE:
        schedule_work(ece_data->shutdown_task);
        break;
    }

    return IRQ_HANDLED;
}
```


Cleanup

- Legacy
 - Turn off the device interrupts
 - `free_irq(pdev, cookie)`
- MSI
 - Turn off device interrupts
 - `pci_disable_msi(pdev)`
- MSI-X
 - Turn off device interrupts and delete array
 - `pci_disable_msix(pdev)`
 - `kfree(entries)`




Top Half / Bottom Half

- Split the activity into quick reaction part and work part
- Top Half
 - Runs in interrupt context, interrupts masked
 - Stashes away info for bottom half
 - Releases HW interrupt context
- Bottom Half
 - Processes received information
 - Sleeps, locks, waits, I/O, etc
 - Re-enables interrupt
 - Workqueue, Tasklet, I/O thread waiting for completion



See example code

- e1000e for interrupt handler
 - Why difference between legacy and MSI handlers?
- i40e for something a bit more advanced...
 - i40e_msix_clean_rings()
 - What is NAPI?!?!
 - i40e_napi_poll()



Workqueue

- Uses a task function, like a callback or handler
- Triggered from other threads, e.g. interrupts
- Runs in full process context – can sleep
- Runs on generic kernel work thread or your own thread
- Use it to "do stuff"



Workqueue Setup

```
struct ixgbe_adapter {  
    ...  
    struct work_struct service_task;  
    ...  
};  
  
static int __devinit ixgbe_probe() {  
    ...  
    INIT_WORK(&adapter->service_task, ixgbe_service_task);  
    ...  
}  
  
static void __dev_exit() {  
    ...  
    cancel_work_sync(&adapter->service_task);  
    ...  
}
```

Workqueue Usage

```
cause = readl(hw_addr + REG_IRQ_CAUSE);
```

```
switch (cause) {
```

Interrupt handler

```
case IXGBE_IRQ_CAUSE_WATCHDOG:
    schedule_work(adapter->service_task);
    break;
}
```

Workqueue code

```
static void ixgbe_service_task(struct work_struct *work)
{
    struct ixgbe_adapter *adapter = container_of(work,
                                                    struct ixgbe_adapter,
                                                    service_task);

    ixgbe_reset_subtask(adapter);
    ixgbe_sfp_detection_subtask(adapter);
    ixgbe_sfp_link_config_subtask(adapter);
    ixgbe_check_overtemp_subtask(adapter);
    ixgbe_watchdog_subtask(adapter);
    ixgbe_fdir_reinit_subtask(adapter);
    ixgbe_check_hang_subtask(adapter);

    ixgbe_service_event_complete(adapter);
}
```

Readings

- LDD3: Chap 10
- ELDD: ppg 72-74, 92-103

